

Restructuring Target Screener

User Guide: How the Agent Screens, Researches and Reports

An AI-assisted workflow, built on Claude Code and the 9fin MCP server, that screens for stressed European corporate credits, researches each candidate, and produces a print-ready tearsheet and covering email in a fraction of the time a manual screen takes.

COVENANT TRIP

Approaching or breaching a financial covenant

MATURITY WALL

Debt maturing in 18–24 months that cannot be refinanced at current leverage

LIQUIDITY SHORTFALL

Insufficient cash and facilities to fund operations through the near term

The three triggers every screen, tearsheet and recommendation is assessed against.

CONTENTS

1. What this agent does	Page 2
2. System architecture	Page 3
3. Folder structure and file naming	Page 4
4. Phase 1: Weekly Screen	Page 5
5. Phase 2: Building the Tearsheet	Page 6
6. Phase 2.5: The Audit Gate	Page 7
7. Phase 3: Delivery	Page 8
8. How to trigger a run	Page 9
9. Continuity across weeks	Page 10
10. Quality bar and limitations	Page 11
11. Quick reference	Page 12

WHAT THIS AGENT DOES

Restructuring teams need to stay on top of newly stressed European corporate credits. Screening for this manually is repetitive: check the screener or the weekly report, look up each name, pull financials and capital structure, judge whether it is a genuine candidate, write it up, and email the team. This workflow automates that process end to end while keeping an analyst in the loop at the two points that matter: deciding which names are worth a full workup, and confirming the email before it goes out.

Everything the agent produces is assessed against the same three triggers shown on the cover page. A recommendation of “further work” always ties back to at least one of them being present or approaching; a “pass” states why none applied.

Two ways to run it

S

SCHEDULED SCREEN

Every Sunday evening, sourced from 9fin's European Weekly Stressed and Distressed Data Report: this week's watchlist entrants plus the top weekly price losers.

M

MANUAL SCREEN

Triggered any time, run against the 9fin bond/loan screener with the confirmed default criteria or custom criteria for that run only.

D


AD HOC DEEP DIVE

Off-cycle Phase 2 workup on a specific name, e.g. following market colour, even where the prior weekly call was “pass”.

SYSTEM ARCHITECTURE

One orchestrating session, three phases, five specialised sub-agents.

Every run starts in a single Claude Code session (the “orchestrator”), which reads **CLAUDE.md** and **workflows/ai-content-agent-workflow.md** for the standing rules and the current step-by-step SOP. The orchestrator does the sequencing and every checkpoint itself; the actual research is delegated to specialised sub-agents, each scoped to exactly the tools it needs.



Sub-agents are scoped in `.claude/agents/*.md`. The three Phase 2 researchers run in parallel per company to save time; the auditor only starts once all three have returned.

FOLDER STRUCTURE AND FILE NAMING

Where everything lives, and how outputs are named so a week's work is easy to find later.

/CLAUDE.md — *standing rules, style guide, the source of truth alongside the workflow SOP*

/workflows/

ai-content-agent-workflow.md — *the step-by-step SOP read at the start of every run*

/.claude/agents/

screeener.md, researcher-capital.md, researcher-financials.md, researcher-news.md, auditor.md

/outputs/ — *generated tearsheets (HTML) and combined PDFs, the durable local record*

/target_list/

target_list.csv — *running record across weeks (see page 10)*

/Templates/

weekly_screen_email_template.md, recipient_names.csv

gold_standard_wagamama_tearsheet.html — *the quality bar (see page 11)*

/resources/{companyname}/ — *uploaded source material for a specific company; outranks 9fin on conflicting facts*

/drafts/ — *work in progress: researcher JSON, audit reports*

/Precedent Emails/ — *archive of what was actually sent*

/docs/ — *this guide, the supporting presentation, and tearsheet screenshots*

File naming convention

STANDARD FORMAT

ddmmyy_companyname_content.filetype

Underscores only, no special characters. For example 280926_merlin_entertainments_tearsheet.html or 270926_weeklyscreen_triage.md. Multiple companies delivered together are combined into one file, e.g.


ddmmyy_weekly_screen.html.

Continuity is the point of the running record

/target_list/target_list.csv is read before every screen and updated after every run. It is what stops the agent re-researching a name from scratch each week; see page 10 for how the carry-forward logic works.

PHASE 1: WEEKLY SCREEN

Two ways in, one shared path from there, one checkpoint at the end.



Why the scheduled run does not use the screener

The bond/loan screener filters on static criteria (rating, sector, size); it has no way to filter on week-over-week price movement, which is the actual signal the Sunday-night run exists to catch. The weekly report already computes that movement natively, so it is the better-fitted source for this cadence. Its two lists (watchlist entrants and top weekly losers) are taken as published, with no additional rating, borrower type, region or status filtering layered on top.

Manual run: default screening criteria

Filter	Default
Status	Priced
Borrower type	Corporate
Region	Europe
Rating (CFR)	BB, B or CCC
Prior week price move	Worse than -2.5pts
Sector, minimum size	No default – full universe unless specified for that run


A one-off set of custom criteria applies to that run only; it does not become the new default unless the analyst says so.

The triage table

For each name: company, one-line business description, key leverage/rating stat, and a two-line recommendation of “further work” or “pass”, stating which of the three triggers drove the call, or why none applied. The run then pauses and waits for the analyst to choose which names move to Phase 2.

PHASE 2: BUILDING THE TEARSHEET

Three researchers work in parallel per company; their combined output becomes one printable page.



Before writing content, the tearsheet is measured against `Templates/gold_standard_wagamama_tearsheet.html` for depth and sourcing rigor, not copied as a content template.

Tearsheet sections


HEADER & OVERVIEW Header block · business description · situation overview · key catalysts	PEOPLE & STRUCTURE Management & sponsor table · group structure chart
CAPITAL & LIQUIDITY Capital structure table · liquidity build · debt holdings	FINANCIALS Revenue splits · summary financials · UFCF/LFCF build and conversion
MARKET & NEWS Annotated trading chart · dated, sourced recent-news table	CALL Triggers to watch · recommendation tied back to the three triggers

Every figure is attributed to a specific source. Conflicting leverage bases are called out rather than blended. A negative FCF conversion is shown as *n.a.* rather than a misleadingly precise negative percentage. Missing data is marked “not available – further diligence required” with the likely source named, never left blank or inferred.

PHASE 2.5: THE AUDIT GATE

A quality check that runs after the three researchers and before a single line of tearsheet HTML is written.

The auditor does no new research. Working only from the three researchers' JSON outputs and anything in `/resources/{companyname}/`, it cross-references every fact stated in more than one place (debt amounts, maturities, EBITDA, leverage, liquidity, covenants) and looks for the same fact reported inconsistently.



This is one clarification round only

If the auditor's second pass still shows a discrepancy as unresolved, the workflow does not loop again and does not pick a value itself. The tearsheet instead shows both conflicting figures against the affected field, e.g. "EUR 810m (9fin) / EUR 840m (Fitch, 06-Jul-26) – unresolved, further diligence required", styled the same way as the FCF "n.a." convention. Where the auditor resolves everything cleanly, whether by the priority rule or a successful clarification round, the tearsheet shows only the correct figure with no visible trace of what was caught.

PHASE 3: DELIVERY

From finished tearsheet HTML to a confirmed email in the analyst's inbox draft or sent folder.

- 1 Build the HTML.** Each tearsheet is A4 portrait. Multiple companies in one week are combined into a single file with a page break between each, named `ddmmyy_weekly_screen.html`; a single company uses its own standard filename.
- 2 Render the PDF.** `render_pdf.js` (Puppeteer, A4, background graphics on) converts the same HTML to `/outputs/ddmmyy_weekly_screen.pdf` – the durable local record of what was sent, independent of the artifact's own lifecycle.
- 3 Publish as a Claude artifact.** Artifacts are private by default, so link sharing (“anyone with the link can view”) is switched on before it goes near the email – otherwise the recipient hits a sign-in wall instead of the report.
- 4 Draft the covering email.** Built from `Templates/weekly_screen_email_template.md`, with the screen criteria, results, recommendation rationale and the artifact link populated in. Greeting rules: one recipient gets a first name, two get both first names, three or more get “Hi team,”. First names are resolved via `Templates/recipient_names.csv`.
- 5 Confirm, every time.** Email content, recipient list, and that the artifact link actually opens are confirmed with the analyst before every send – a standing constraint, not a one-off setup step.
- 6 Send via Gmail.** The email carries the artifact link in the body only; no PDF attachment.

HANDLE THE ARTIFACT LINK LIKE THE PDF

A link trades the PDF's total containment (only the recipient ever had the file) for reach-by-link (anyone who gets or guesses the URL can open it). The tearsheet contains licensed 9fin data and rating-agency content, so the link is treated the same way as the PDF: internal use only, not for onward forwarding or posting.

HOW TO TRIGGER A RUN

What to say to Claude Code to start each type of run.

Scheduled Sunday-night screen

"Run this week's screen."

No clarifying questions are asked for this path. The agent goes straight to the weekly report, builds the triage table, and pauses for your input on which names to progress.

Manual screen with default criteria

"Run an ad hoc screen with the default criteria."

Manual screen with custom criteria

"Screen for German or French corporates, B or CCC rated, with 2027–2028 maturities."

Any field left unspecified falls back to the confirmed default. This one-off criteria set applies only to that run unless you say to make it the new default.

Ad hoc deep dive on a specific name

"Do a full Phase 2 workup on [Company], even though we passed on it last month – there's fresh market colour on a maturity extension."

Changing or extending the workflow itself

"I want to add a fourth trigger for asset-sale-dependent liquidity plans."

Build/change requests are handled differently from a routine run: at least three clarifying questions are asked before starting, and a plan is presented for approval before any multi-step change is executed.

Where it always pauses

Checkpoint	What it waits for
After the triage table (Phase 1)	Which names to take forward to Phase 2
Before sending (Phase 3)	Confirmation of email content, recipients, and that the artifact link opens

Anything that does not fit the existing scope in a given week (a new data situation, a screener error, an ambiguous instruction) is also a stop-and-ask point, rather than being improvised around.

CONTINUITY ACROSS WEEKS

The screen does not start from zero every Sunday.

/target_list/target_list.csv is read at the start of every run, before 9fin is queried. For any name that reappears on the current week's screen:

- Its prior recommendation and the date it was last assessed are noted first.
- The full workup is only redone if something material has changed: a price move, a rating action, news, or new financials.
- Otherwise, the prior view is carried forward with a note such as “no material change since 26-Sep-2026”.

At the end of every run, the target list is updated with that week's names, recommendations and status, so the next run starts from a current record.

Example row from the running record

Company	Sector	Net Leverage	Next Maturity	Recommendation	Date Assessed
Ahlstrom-Munksjo	Paper & Forest Products	5.2x	EUR 610m SSNs, Feb-2028	Watchlist	26-Sep-2026

Notes field for this row: unrefinanceable maturity wall (EUR 610m SSNs Feb-2028) with an RCF springing clause in Nov-2027; elevated leverage at 5.2x IFRS 16; TLB extended to 2030 but SSN refinancing outstanding.

Why this matters

A weekly screen that reprocessed every recurring name in full would burn time re-deriving conclusions that have not changed. Carrying forward a “no material change” view keeps the analyst's attention on what actually moved this week, while still surfacing the full prior context (sector, sponsor, leverage, liquidity, next maturity, rating, and the last recommendation) for anything that reappears.

What “good” looks like, and where the workflow's edges are.

Wagamama UK pan-Asian casual dining – Wagamama Group Limited (E330m 8.5% SSNs 2030)		Prepared for internal use – CMCK Associates Last updated: 27-Sep-2026 Coverage: B2 (Moody's, neg.) / B (S&P, neg.)	
SECTOR / INDUSTRY Consumer Discretionary – Restaurants OWNERSHIP / SPONSOR Apollo Global Management (via THG)	COUNTRY United Kingdom (86% of revenue) LAST UPDATE DATE 27-Sep-2026	COMPANY TYPE Private, sponsor-owned bond issuer COVERAGE STATUS 96n covered, S&P report reviewed 26-Jun-26	TICKER n/a – undisted RESTRICTED GROUP Wagamama Ltd / Wagamama (Redding) Ltd
WATCHLIST – MONITOR No trigger currently active. Liquidity is the primary line concern (conditional S&P "adequate" rating, negative FOCF through FY26, and the RCP's upcoming coverage is approaching activation 27.4% drawn vs. 40% threshold), – both warrant quarterly monitoring. No near-term maturity wall (SSNs due Jun-2030).			

Page 1 – header, business description, situation overview, key catalysts

Page 2 – management, capital structure, liquidity build, group structure

Page 3 – debt holdings, revenue splits, summary financials, FCF build

Page 4 – trading chart, recent news, recommendation

Restructuring Target Screener · User Guide

QUICK REFERENCE

A one-page cheat sheet for the routine week.

You want to...	Say...	You get...
Run the weekly screen	"Run this week's screen"	Triage table, then a pause for your input
Run an ad hoc screen	State the criteria, or ask for the default set	Triage table from the bond/loan screener
Progress names to Phase 2	Name the companies from the triage table	Parallel research, audit, then tearsheet(s)
Investigate a "pass" name	Describe the new colour and name the company	A full Phase 2 workup, off-cycle
Send the week's results	Approve the drafted email	Sent email with the shared artifact link
Change how the workflow behaves	Describe the change you want	Clarifying questions, then a plan to approve

The three triggers, restated

COVENANT TRIP

Approaching or breaching a financial covenant

MATURITY WALL

18–24 month maturities that cannot be refinanced at current leverage

LIQUIDITY SHORTFALL

Insufficient cash and facilities for the near term

Repo map

/CLAUDE.md · /workflows · /.claude/agents · /outputs · /target_list · /Templates · /resources/{companyname} · /drafts · /Precedent Emails · /docs

Further reading in this repository

- README.md – project overview and demo videos
- docs/Presentation/9fin MCP Agent Supporting Materials.pdf – the fuller walkthrough and thinking behind the workflow
- workflows/ai-content-agent-workflow.md and CLAUDE.md – the source-of-truth SOP and standing rules